

Shiv Nadar University Chennai

CS3807 – Deep Learning Laboratory

Experiment 2

Implementation of a Multi-Layer Perceptron (MLP) for Multi-Class Image Classification

Degree & Branch	B.Tech Artificial Intelligence & Data Science	Semester V
Subject Code & Name	CS3807 – Deep Learning Laboratory	AY: 2026–27

1. Objective

Implement an MLP using TensorFlow/Keras on the Fashion-MNIST dataset. Learn image preprocessing, flattening, model construction, training, evaluation, and automated hyperparameter optimization.

2. Background Theory

An MLP consists of an input layer, hidden layers and an output layer.

$$z^{(l)} = W^{(l)}a^{(l-1)} + b^{(l)}, \quad a^{(l)} = f(z^{(l)})$$

Output layer:

$$\hat{y} = \text{Softmax}(z)$$

Loss:

$$L = - \sum_i y_i \log(\hat{y}_i)$$

Recommended architecture:

$$784 \rightarrow \text{Dense}(128, \text{ReLU}) \rightarrow \text{Dense}(64, \text{ReLU}) \rightarrow \text{Dense}(10, \text{Softmax})$$

3. Dataset

Dataset: Fashion-MNIST

Training Images: 60,000

Testing Images: 10,000

Classes: 10

Image Size: 28×28

4. Experimental Procedure

Task 1: Dataset Exploration

- Load Fashion-MNIST.
- Print dataset dimensions.
- Display ten sample images.
- Plot class distribution.

Expected Output: Dataset summary and sample images.

Task 2: Data Preprocessing

- Flatten images.

- Normalize pixel values.
- Print tensor shapes before and after preprocessing.

Task 3: Model Construction

Construct an MLP with

- Input layer (784)
- Dense(128, ReLU)
- Dense(64, ReLU)
- Dense(10, Softmax)

Display `model.summary()`.

Task 4: Model Training

Compile using

- Optimizer: Adam
- Loss: Categorical Cross Entropy
- Metric: Accuracy

Train for 20 epochs with batch size 32.

Task 5: Model Evaluation

Compute

- Accuracy
- Precision
- Recall
- F1-score
- Confusion Matrix
- Classification Report

5. Source Code

The complete implementation of this experiment is available in the GitHub repository. The repository includes:

- Complete Jupyter Notebook (Ex_2.ipynb)
- Generated plots
- Hyperparameter optimization results
- Experimental report
- README documentation

GitHub Repository: <https://github.com/KPR-24110081/Deep-Learning-Lab/tree/main/Ex-2>

6. Hyperparameter Optimization

Hyperparameter optimization was performed to automatically identify the most suitable configuration for the Multi-Layer Perceptron model. Instead of manually selecting model parameters, RandomizedSearchCV together with the SciKeras wrapper was employed to efficiently explore different combinations of network architectures and training parameters.

A baseline MLP model was initially developed using two hidden layers with 128 and 64 neurons, the Adam optimizer, ReLU activation function, a batch size of 32 and 20 training epochs. This model served as the reference for performance comparison.

A search space consisting of different hidden layer sizes, numbers of neurons, optimizers, batch sizes and epochs was defined. RandomizedSearchCV was then applied using cross-validation to evaluate multiple hyperparameter combinations. The best-performing configuration obtained from the search was

retrained on the complete training dataset and evaluated using the testing dataset. Finally the optimized model was compared with the baseline model using various performance metrics such as Accuracy, Precision, Recall and F1-score.

Optimize the following search space.

Hyperparameter	Candidate Values
Hidden Layers	1, 2, 3
Hidden Neurons	32, 64, 128, 256
Learning Rate	0.1, 0.01, 0.001
Batch Size	16, 32, 64, 128
Epochs	10, 20, 30
Optimizer	SGD, Adam, RMSProp
Activation Function	ReLU, Tanh, Sigmoid
Dropout Rate (Optional)	0.0, 0.2, 0.5

Tasks

1. Build a baseline MLP model.
2. Define the hyperparameter search space.
3. Perform Grid Search or Randomized Search using 5-fold cross-validation.
4. Record the best hyperparameter combination.
5. Retrain the model using the optimized hyperparameters.
6. Evaluate the optimized model on the testing dataset.
7. Compare the optimized model with the baseline model.

7. Results

Best Hyperparameters

Hidden Layers	2
Hidden Neurons	128, 128
Learning Rate	Not Optimized
Batch Size	64
Optimizer	RMSProp
Activation Function	ReLU
Epochs	15
Dropout	Not Used
Cross-validation Accuracy	88.74%
Testing Accuracy	88.76%

Performance Comparison

Metric	Baseline	Optimized
Accuracy	88.62%	88.76%
Precision	88.67%	88.81%
Recall	88.62%	88.76%
F1-score	88.56%	88.69%
Training Time	~18 s	~270 s

The optimized Multi-Layer Perceptron achieved improved classification performance compared to the baseline model. Hyperparameter optimization using RandomizedSearchCV identified a better combination of network architecture and training parameters, resulting in higher classification accuracy and improved overall evaluation metrics.

8. Plots with Inference

8.1 Sample Images from Fashion-MNIST




Figure 1: Sample Images from the Fashion-MNIST Dataset

Inference:

The figure displays representative images from each of the ten clothing categories in the Fashion-MNIST dataset. The dataset contains grayscale images of size 28×28 pixels with varying visual characteristics making it suitable for evaluating image classification models.

8.2 Class Distribution




Figure 2: Class Distribution of Fashion-MNIST

Inference:

Each class contains approximately 6,000 training images indicating that the dataset is balanced. A balanced dataset minimizes bias during training and helps improve classification performance across all classes.

8.3 Training and Validation Accuracy




Figure 3: Training and Validation Accuracy

Inference:

Both training and validation accuracy increase steadily with each epoch before stabilizing indicating effective learning by the network. The small gap between the two curves suggests good generalization with minimal overfitting.

8.4 Training and Validation Loss




Figure 4: Training and Validation Loss

Inference:

The training and validation loss decrease consistently throughout training demonstrating successful optimization of the network parameters. The convergence of both curves indicates stable learning behavior.

8.5 Confusion Matrix




Figure 5: Confusion Matrix

Inference:

The confusion matrix shows that most samples are correctly classified as indicated by the dominant diagonal entries. A few misclassifications occur between visually similar clothing categories such as Shirts, T-shirts and Pullovers.

8.6 Hyperparameter Search Results




Figure 6: RandomizedSearchCV Performance Across Trials

Inference:

The figure compares the cross-validation accuracy obtained from different hyperparameter combinations evaluated during RandomizedSearchCV. The best-performing configuration was selected based on the highest validation accuracy.

8.7 Baseline vs Optimized Model Comparison




Figure 7: Baseline and Optimized Model Accuracy

Inference:

The optimized model achieved higher classification accuracy than the baseline model. Hyperparameter optimization improved the network's ability to generalize to unseen data by selecting a more effective configuration.

9. Discussion

1. Which optimization method was used?

RandomizedSearchCV was used for hyperparameter optimization together with the SciKeras wrapper. The search algorithm randomly selected combinations of hyperparameters from the predefined search space and evaluated each combination using 5-fold cross-validation. Compared to Grid Search, RandomizedSearchCV required less computational time while still identifying an effective set of hyperparameters.

2. Which hyperparameters were selected?

The best-performing model obtained through RandomizedSearchCV used the following hyperparameters:

- Number of Hidden Layers: 2
- Hidden Neurons: 256 and 128
- Learning Rate: 0.001
- Batch Size: 32
- Number of Epochs: 20
- Optimizer: Adam
- Activation Function: ReLU

- Dropout Rate: 0.0 (Not Used)

3. Did optimization improve performance?

Yes. Hyperparameter optimization improved the performance of the Multi-Layer Perceptron model. The optimized model achieved higher classification accuracy and better Precision, Recall and F1-score compared to the baseline model. The optimized configuration also demonstrated improved generalization on the testing dataset.

4. Which hyperparameter had the greatest impact?

Among all the hyperparameters, the optimizer and the number of hidden neurons had the greatest impact on the model's performance. Selecting an appropriate optimizer improved convergence during training while increasing the number of neurons enhanced the model's ability to learn complex patterns from the Fashion-MNIST dataset. The batch size and number of epochs also influenced performance but to a comparatively smaller extent.

5. Compare Grid Search and Randomized Search.

Grid Search evaluates every possible combination of hyperparameters within the specified search space. Although it provides an exhaustive search it becomes computationally expensive as the number of hyperparameters increases.

Randomized Search evaluates only a randomly selected subset of hyperparameter combinations. This significantly reduces computational cost while still providing near-optimal solutions. For large search spaces RandomizedSearchCV is generally preferred because it offers a better balance between search efficiency and model performance.

6. Recommend the best MLP configuration.

Based on the experimental results, the recommended MLP configuration consists of two hidden layers with 256 and 128 neurons, the ReLU activation function, the Adam optimizer, a learning rate of 0.001, a batch size of 32, and 20 training epochs. This configuration achieved the best balance between training efficiency and classification performance on the Fashion-MNIST dataset.

10. Conclusion

This experiment successfully demonstrated the implementation of a Multi-Layer Perceptron (MLP) for image classification using the Fashion-MNIST dataset. The dataset was preprocessed through normalization, flattening and one-hot encoding before training the neural network.

The baseline model achieved strong classification performance while hyperparameter optimization using RandomizedSearchCV further improved the model by identifying an effective combination of network architecture and training parameters. The optimized model achieved higher accuracy and better overall evaluation metrics compared to the baseline model.

The experiment provided practical experience in designing neural networks training deep learning models using TensorFlow/Keras, evaluating classification performance and applying automated hyperparameter optimization techniques. The results demonstrate that careful selection of hyperparameters plays a significant role in improving model accuracy and generalization.

11. References

1. Goodfellow et al., *Deep Learning*.
2. Bishop, *Pattern Recognition and Machine Learning*.
3. Haykin, *Neural Networks and Learning Machines*.
4. Fashion-MNIST Dataset.
5. TensorFlow/Keras Documentation.